

Regressão Linear - TE2

April 29, 2025

```
[2]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[3]: casas = pd.read_csv('USA_Housing.csv')
```

```
[4]: casas.head()
```

```
[4]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	\
0	79545.458574	5.682861	7.009188	
1	79248.642455	6.002900	6.730821	
2	61287.067179	5.865890	8.512727	
3	63345.240046	7.188236	5.586729	
4	59982.197226	5.040555	7.839388	

	Avg. Area Number of Bedrooms	Area Population	Price	\
0	4.09	23086.800503	1.059034e+06	
1	3.09	40173.072174	1.505891e+06	
2	5.13	36882.159400	1.058988e+06	
3	3.26	34310.242831	1.260617e+06	
4	4.23	26354.109472	6.309435e+05	

	Address
0	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	9127 Elizabeth Stravenue\nDanielstown, WI 06482...
3	USS Barnett\nFPO AP 44820
4	USNS Raymond\nFPO AE 09386

0.0.1 Modelo supervisionado

- separação entre base de treino e base de teste

0.0.2 Modelo não supervisionado

- não existe a separação entre base de treino e teste

```
[5]: casas.info()
```

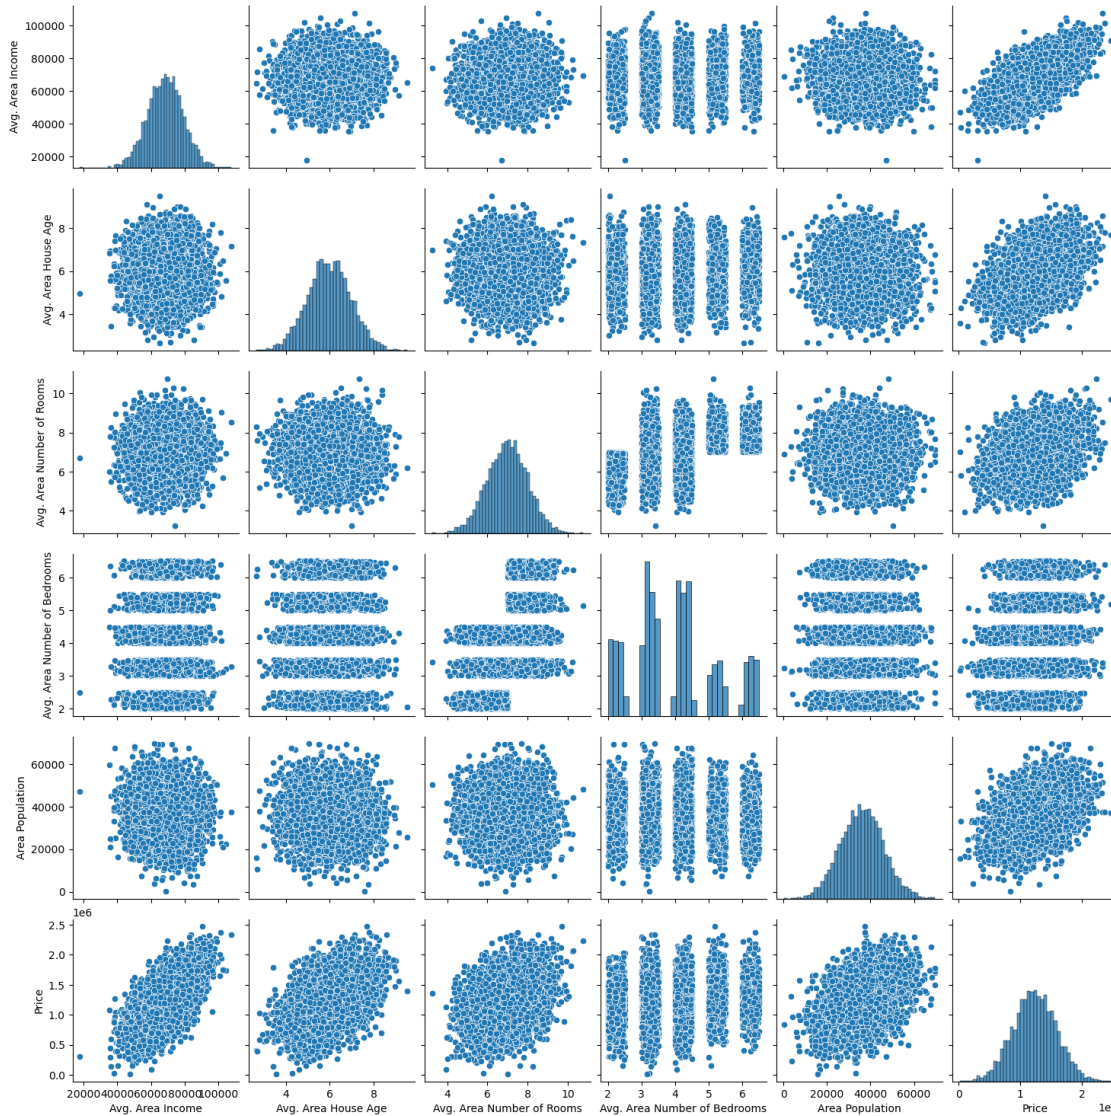
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                      5000 non-null   float64
1   Avg. Area House Age                   5000 non-null   float64
2   Avg. Area Number of Rooms             5000 non-null   float64
3   Avg. Area Number of Bedrooms          5000 non-null   float64
4   Area Population                       5000 non-null   float64
5   Price                                 5000 non-null   float64
6   Address                               5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB

```

```
[6]: sns.pairplot(casas)
```

```
[6]: <seaborn.axisgrid.PairGrid at 0x7fa925e78c50>
```



```
[7]: casas.corr()
```

/tmp/ipykernel_253299/2249505781.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
casas.corr()
```

```
[7]:
```

	Avg. Area Income	Avg. Area House Age	\
Avg. Area Income	1.000000	-0.002007	
Avg. Area House Age	-0.002007	1.000000	
Avg. Area Number of Rooms	-0.011032	-0.009428	
Avg. Area Number of Bedrooms	0.019788	0.006149	

Area Population	-0.016234	-0.018743
Price	0.639734	0.452543

	Avg. Area Number of Rooms \
Avg. Area Income	-0.011032
Avg. Area House Age	-0.009428
Avg. Area Number of Rooms	1.000000
Avg. Area Number of Bedrooms	0.462695
Area Population	0.002040
Price	0.335664

	Avg. Area Number of Bedrooms	Area Population \
Avg. Area Income	0.019788	-0.016234
Avg. Area House Age	0.006149	-0.018743
Avg. Area Number of Rooms	0.462695	0.002040
Avg. Area Number of Bedrooms	1.000000	-0.022168
Area Population	-0.022168	1.000000
Price	0.171071	0.408556


	Price
Avg. Area Income	0.639734
Avg. Area House Age	0.452543
Avg. Area Number of Rooms	0.335664
Avg. Area Number of Bedrooms	0.171071
Area Population	0.408556
Price	1.000000

```
[8]: sns.heatmap(casas.corr())
```

```
/tmp/ipykernel_253299/254642754.py:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it will
default to False. Select only valid columns or specify the value of numeric_only
to silence this warning.
```

```
sns.heatmap(casas.corr())
```

```
[8]: <AxesSubplot: >
```



```
[9]: casas.columns
```

```
[9]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
        'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],
        dtype='object')
```

```
[10]: casas[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
            'Avg. Area Number of Bedrooms', 'Area Population', 'Price']].corr()
```

```
[10]:
```

	Avg. Area Income	Avg. Area House Age	\
Avg. Area Income	1.000000	-0.002007	
Avg. Area House Age	-0.002007	1.000000	
Avg. Area Number of Rooms	-0.011032	-0.009428	
Avg. Area Number of Bedrooms	0.019788	0.006149	
Area Population	-0.016234	-0.018743	
Price	0.639734	0.452543	


	Avg. Area Number of Rooms \
Avg. Area Income	-0.011032
Avg. Area House Age	-0.009428
Avg. Area Number of Rooms	1.000000
Avg. Area Number of Bedrooms	0.462695
Area Population	0.002040
Price	0.335664

	Avg. Area Number of Bedrooms	Area Population \
Avg. Area Income	0.019788	-0.016234
Avg. Area House Age	0.006149	-0.018743
Avg. Area Number of Rooms	0.462695	0.002040
Avg. Area Number of Bedrooms	1.000000	-0.022168
Area Population	-0.022168	1.000000
Price	0.171071	0.408556

	Price
Avg. Area Income	0.639734
Avg. Area House Age	0.452543
Avg. Area Number of Rooms	0.335664
Avg. Area Number of Bedrooms	0.171071
Area Population	0.408556
Price	1.000000


```
[12]: sns.heatmap(casas[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number
    of Rooms',
    'Avg. Area Number of Bedrooms', 'Area Population', 'Price']].corr(),
    annot=True)
```

```
[12]: <AxesSubplot: >
```



```
[13]: sns.histplot(casas['Price'])
```

```
[13]: <AxesSubplot: xlabel='Price', ylabel='Count'>
```



Modelo de regressão

- X -> VARIÁVEIS PREDITORAS
- Y -> VARIÁVEL QUE QUERO PREDIZER (DESCOBRIR) - ALVO

[14]: `casas.info()`

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 5000 entries, 0 to 4999

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	Avg. Area Income	5000 non-null	float64
1	Avg. Area House Age	5000 non-null	float64
2	Avg. Area Number of Rooms	5000 non-null	float64
3	Avg. Area Number of Bedrooms	5000 non-null	float64
4	Area Population	5000 non-null	float64
5	Price	5000 non-null	float64
6	Address	5000 non-null	object

dtypes: float64(6), object(1)

memory usage: 273.6+ KB


```
[15]: casas.columns
```

```
[15]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
        'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],  
        dtype='object')
```

```
[16]: X = casas[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
        'Avg. Area Number of Bedrooms', 'Area Population']]  
Y = casas['Price']
```

```
[17]: X
```

```
[17]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	\
0	79545.458574	5.682861	7.009188	
1	79248.642455	6.002900	6.730821	
2	61287.067179	5.865890	8.512727	
3	63345.240046	7.188236	5.586729	
4	59982.197226	5.040555	7.839388	
...	
4995	60567.944140	7.830362	6.137356	
4996	78491.275435	6.999135	6.576763	
4997	63390.686886	7.250591	4.805081	
4998	68001.331235	5.534388	7.130144	
4999	65510.581804	5.992305	6.792336	

	Avg. Area Number of Bedrooms	Area Population
0	4.09	23086.800503
1	3.09	40173.072174
2	5.13	36882.159400
3	3.26	34310.242831
4	4.23	26354.109472
...
4995	3.46	22837.361035
4996	4.02	25616.115489
4997	2.13	33266.145490
4998	5.44	42625.620156
4999	4.07	46501.283803

```
[5000 rows x 5 columns]
```

```
[18]: Y
```

```
[18]:
```

0	1.059034e+06
1	1.505891e+06
2	1.058988e+06
3	1.260617e+06

```

4          6.309435e+05
...
4995      1.060194e+06
4996      1.482618e+06
4997      1.030730e+06
4998      1.198657e+06
4999      1.298950e+06
Name: Price, Length: 5000, dtype: float64

```

0.0.3 Divisão do conjunto de dados

0.0.4 %pip install scikit-learn

```
[19]: from sklearn.model_selection import train_test_split
```

```
[20]: X_treino, X_teste, Y_treino, Y_teste = train_test_split(X,Y,
                                                             train_size=0.7,
                                                             random_state=46)
```

```
[21]: X_treino
```

```
[21]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms \
456	62653.092462	6.543984	7.884326
4480	71481.034926	5.188492	7.152361
3564	63889.411593	5.548089	6.357831
587	51918.546873	5.892497	6.708809
2634	42814.993038	5.247613	6.080981
...
3933	71229.357964	4.850191	6.978128
3787	50362.538095	5.582574	4.608843
658	73829.777741	6.130263	6.182843
2451	63421.903955	7.594954	8.777735
2490	65010.553474	6.299952	7.529846

	Avg. Area Number of Bedrooms	Area Population
456	3.28	41467.867658
4480	4.32	45246.174103
3564	2.11	34442.099648
587	3.37	24734.406520
2634	4.02	41426.389765
...
3933	4.15	27168.842194
3787	2.40	23611.056225
658	3.50	51385.523241
2451	3.09	11511.387050
2490	6.32	34706.866627

```
[3500 rows x 5 columns]
```

```
[22]: Y_treino
```

```
[22]: 456      1.382110e+06
      4480    1.530013e+06
      3564    1.134126e+06
      587     6.610434e+05
      2634    4.525302e+05
      ...
      3933    6.997873e+05
      3787    3.141678e+05
      658     1.376493e+06
      2451    1.432318e+06
      2490    1.305186e+06
      Name: Price, Length: 3500, dtype: float64
```

```
[23]: X_treino.shape[0]
```

```
[23]: 3500
```

```
[24]: X_teste.shape[0]
```

```
[24]: 1500
```

```
[25]: Y_treino.shape
```

```
[25]: (3500,)
```

```
[26]: Y_teste.shape
```

```
[26]: (1500,)
```

Criando o modelo de treino

```
[27]: from sklearn.linear_model import LinearRegression
```

```
[28]: LR=LinearRegression()
```

```
[29]: LR.fit(X_treino,Y_treino )
```

```
[29]: LinearRegression()
```


```
[30]: y_predicao = LR.predict(X_teste)
```

```
[31]: y_predicao
```

```
[31]: array([1176591.88192438,  997847.03239798, 1179674.83216237, ...,
        934885.47457475, 1497530.06928932,  866365.31013406])
```

```
[32]: plt.scatter(Y_teste, y_predicao)
```

[32]: <matplotlib.collections.PathCollection at 0x7fa913c68590>



```
[33]: from sklearn import metrics
```

```
[34]: metrics.mean_absolute_error(Y_teste, y_predicao)
```

[34]: 79806.83254169731

Y_teste

y_predicao

$((Y_teste - y_predicao)^2)1/2$

```
[35]: metrics.mean_squared_error(Y_teste, y_predicao)
```

[35]: 9993926845.480902

Y_teste

y_predicao

$((Y_teste - y_predicao)^2)$

Atividade

Com base na regressão linear vista em aula, remova um dos preditores e compare o resultado obtido com esse modelo.

```
[37]: X2 = casas[['Avg. Area Income', 'Avg. Area House Age', 'Area Population']]  
      Y2 = casas['Price']
```

```
[38]: X2_treino, X2_teste, Y2_treino, Y2_teste = train_test_split(X2,Y2,  
                                                                train_size=0.7,  
                                                                random_state=46)
```

```
[39]: LR2=LinearRegression()
```


```
[40]: LR2.fit(X2_treino,Y2_treino)
```

```
[40]: LinearRegression()
```

```
[41]: Y2_predicao = LR2.predict(X2_teste)
```

```
[42]: plt.scatter(Y2_teste, Y2_predicao)
```

```
[42]: <matplotlib.collections.PathCollection at 0x7fa911a6d750>
```



```
[43]: metrics.mean_absolute_error(Y2_teste, Y2_predicao)
```

```
[43]: 127067.82406794395
```

```
[ ]:
```