

Regressão_Linear

April 29, 2025

0.1 Regressão Linear

- **Regressão Linear Simples:** um atributo independente e um atributo dependente.
- **Regressão Linear Univariada:** um atributo dependente e um conjunto de atributos independentes
- **Regressão Multivariada:** mais de um atributo dependente.

0.1.1 Regressão Linear Simples

- Y - variável dependente
- X - variável independente
- β_0 - a interceptação
- β_1 - a inclinação

Equação da reta para a RLS

$$e = \beta_0 + \beta_1 * X$$

```
[3]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
```

```
[4]: # Valores de para x de 0 até 19, totalizando 20 pontos
x = np.arange(0,19)
# Matriz para multiplicação dos dados
A=np.array([x, np.ones(19)])
# conjunto de alvos para regressão
y = [22, 24, 23, 20, 23, 30, 22, 24, 24, 20, 21, 30, 24, 20, 28, 20, 22, 24, 27]
```

```
[5]: x
```

```
[5]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
        17, 18])
```

```
[6]: print(A)
```

```
[[ 0.  1.  2.  3.  4.  5.  6.  7.  8.  9. 10. 11. 12. 13. 14. 15. 16. 17.
 18.]
 [ 1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.
 1.]]
```

```
[7]: w=np.linalg.lstsq(A.T,y)[0]
```

/tmp/ipykernel_2595866/1867454896.py:1: FutureWarning: `rcond` parameter will change to the default of machine precision times ``max(M, N)`` where M and N are the input matrix dimensions.


To use the future default and silence this warning we advise to pass `rcond=None`, to keep using the old, explicitly pass `rcond=-1`.

```
w=np.linalg.lstsq(A.T,y)[0]
```

```
[8]: # na posição [0] temos o o valor de beta 1 e na posição 1 o valor de beta 0
print(w)
```

```
[ 0.06666667 22.97894737]
```

```
[9]: linha = w[0]*x+w[1]
plt.plot(x, linha, 'm-')
plt.plot(x, y, 'b')
plt.show()
```



```
[10]: p19=w[0]*19+w[1]
p19
```

```
[10]: 24.2456140350877
```

```
[ ]:
```

```
[11]: import seaborn as sns
```

```
[12]: casas = pd.read_csv('USA_Housing.csv')
```

0.2 Análise estatística do quadro de dados

```
[13]: casas.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                      5000 non-null   float64
1   Avg. Area House Age                   5000 non-null   float64
2   Avg. Area Number of Rooms             5000 non-null   float64
3   Avg. Area Number of Bedrooms          5000 non-null   float64
4   Area Population                       5000 non-null   float64
5   Price                                 5000 non-null   float64
6   Address                               5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

```
[14]: casas.head()
```

```
[14]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	\
0	79545.458574	5.682861	7.009188	
1	79248.642455	6.002900	6.730821	
2	61287.067179	5.865890	8.512727	
3	63345.240046	7.188236	5.586729	
4	59982.197226	5.040555	7.839388	

	Avg. Area Number of Bedrooms	Area Population	Price	\
0	4.09	23086.800503	1.059034e+06	
1	3.09	40173.072174	1.505891e+06	
2	5.13	36882.159400	1.058988e+06	
3	3.26	34310.242831	1.260617e+06	
4	4.23	26354.109472	6.309435e+05	

	Address
0	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	9127 Elizabeth Stravenue\nDanielstown, WI 06482...
3	USS Barnett\nFPO AP 44820
4	USNS Raymond\nFPO AE 09386

```
[15]: casas.tail()
```

```
[15]:      Avg. Area Income  Avg. Area House Age  Avg. Area Number of Rooms  \
4995      60567.944140                7.830362                6.137356
4996      78491.275435                6.999135                6.576763
4997      63390.686886                7.250591                4.805081
4998      68001.331235                5.534388                7.130144
4999      65510.581804                5.992305                6.792336

      Avg. Area Number of Bedrooms  Area Population      Price  \
4995                3.46      22837.361035  1.060194e+06
4996                4.02      25616.115489  1.482618e+06
4997                2.13      33266.145490  1.030730e+06
4998                5.44      42625.620156  1.198657e+06
4999                4.07      46501.283803  1.298950e+06

                                Address
4995      USNS Williams\nFPO AP 30153-7653
4996      PSC 9258, Box 8489\nAPO AA 42991-3352
4997  4215 Tracy Garden Suite 076\nJoshualand, VA 01...
4998      USS Wallace\nFPO AE 73316
4999  37778 George Ridges Apt. 509\nEast Holly, NV 2...
```

```
[16]: casas.describe().transpose()
```

```
[16]:      count      mean      std  \
Avg. Area Income    5000.0  6.858311e+04  10657.991214
Avg. Area House Age    5000.0  5.977222e+00   0.991456
Avg. Area Number of Rooms    5000.0  6.987792e+00   1.005833
Avg. Area Number of Bedrooms    5000.0  3.981330e+00   1.234137
Area Population    5000.0  3.616352e+04   9925.650114
Price    5000.0  1.232073e+06  353117.626581

      min      25%      50%  \
Avg. Area Income  17796.631190  61480.562388  6.880429e+04
Avg. Area House Age    2.644304    5.322283  5.970429e+00
Avg. Area Number of Rooms    3.236194    6.299250  7.002902e+00
Avg. Area Number of Bedrooms    2.000000    3.140000  4.050000e+00
Area Population    172.610686  29403.928702  3.619941e+04
Price    15938.657923  997577.135049  1.232669e+06

      75%      max
Avg. Area Income  7.578334e+04  1.077017e+05
Avg. Area House Age    6.650808e+00  9.519088e+00
Avg. Area Number of Rooms    7.665871e+00  1.075959e+01
Avg. Area Number of Bedrooms    4.490000e+00  6.500000e+00
Area Population    4.286129e+04  6.962171e+04
```


Price 1.471210e+06 2.469066e+06

0.3 Análise Bivariada dos dados

- Boxplot
- gráfico de barras
- Pairplot
- Heatmap correlação dos dados

```
[17]: sns.boxplot(casas)
```

```
[17]: <AxesSubplot: >
```



```
[18]: Q1=casas.quantile(0.25, numeric_only=True)
      Q3=casas.quantile(0.75, numeric_only=True)
      IRQ = Q3-Q1
```

```
[19]: print(Q1)
```

Avg. Area Income	61480.562388
Avg. Area House Age	5.322283
Avg. Area Number of Rooms	6.299250

```

Avg. Area Number of Bedrooms      3.140000
Area Population                   29403.928702
Price                             997577.135049
Name: 0.25, dtype: float64

```

```
[20]: print(Q3)
```

```

Avg. Area Income      7.578334e+04
Avg. Area House Age    6.650808e+00
Avg. Area Number of Rooms  7.665871e+00
Avg. Area Number of Bedrooms  4.490000e+00
Area Population         4.286129e+04
Price                  1.471210e+06
Name: 0.75, dtype: float64

```

```
[21]: print(IRQ)
```

```

Avg. Area Income      14302.776278
Avg. Area House Age    1.328525
Avg. Area Number of Rooms  1.366621
Avg. Area Number of Bedrooms  1.350000
Area Population         13457.362067
Price                  473633.069163
dtype: float64

```

```
[22]: contador = casas[(casas<(Q1-1.5*IRQ)) | (casas > (Q3+1.5*IRQ))].count()

df_contagem = pd.DataFrame(contador, columns=['contagem de outliers'])
```

```

/tmp/ipykernel_2595866/1575016419.py:1: FutureWarning: Automatic reindexing on
DataFrame vs Series comparisons is deprecated and will raise ValueError in a
future version. Do `left, right = left.align(right, axis=1, copy=False)` before
e.g. `left == right`
    contador = casas[(casas<(Q1-1.5*IRQ)) | (casas > (Q3+1.5*IRQ))].count()

```

```
[23]: df_contagem
```

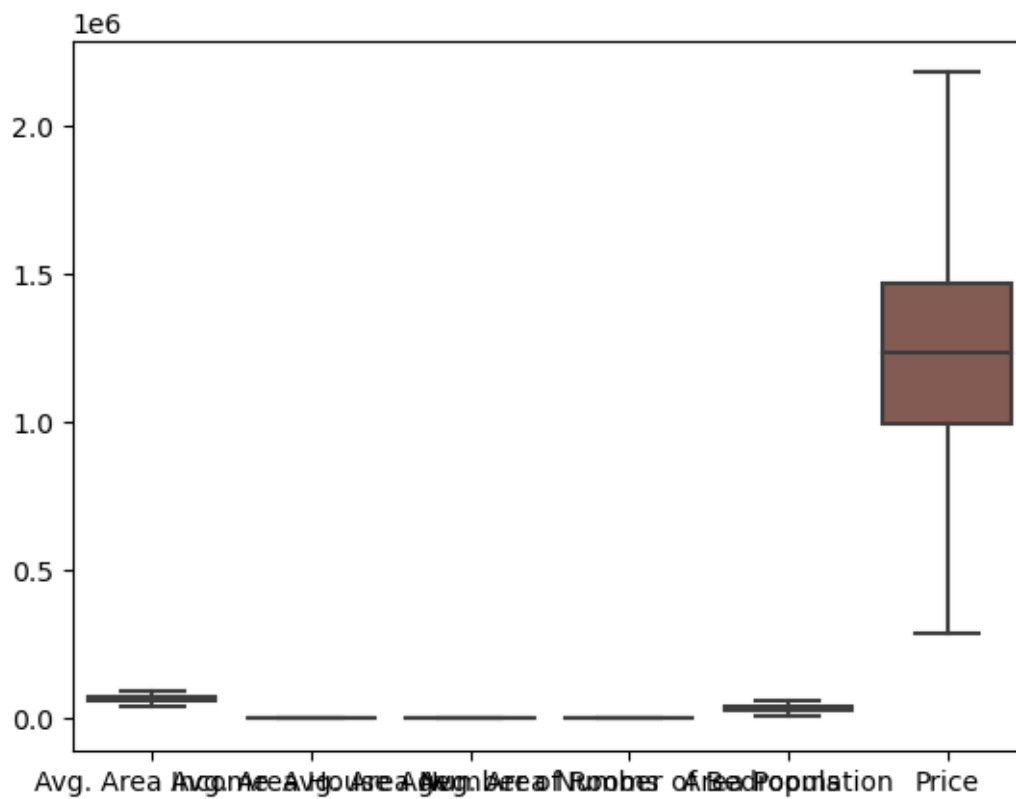
```

[23]:                                     contagem de outliers
Avg. Area Income      32
Avg. Area House Age    25
Avg. Area Number of Rooms  24
Avg. Area Number of Bedrooms  0
Area Population         30
Price                  35
Address                0

```

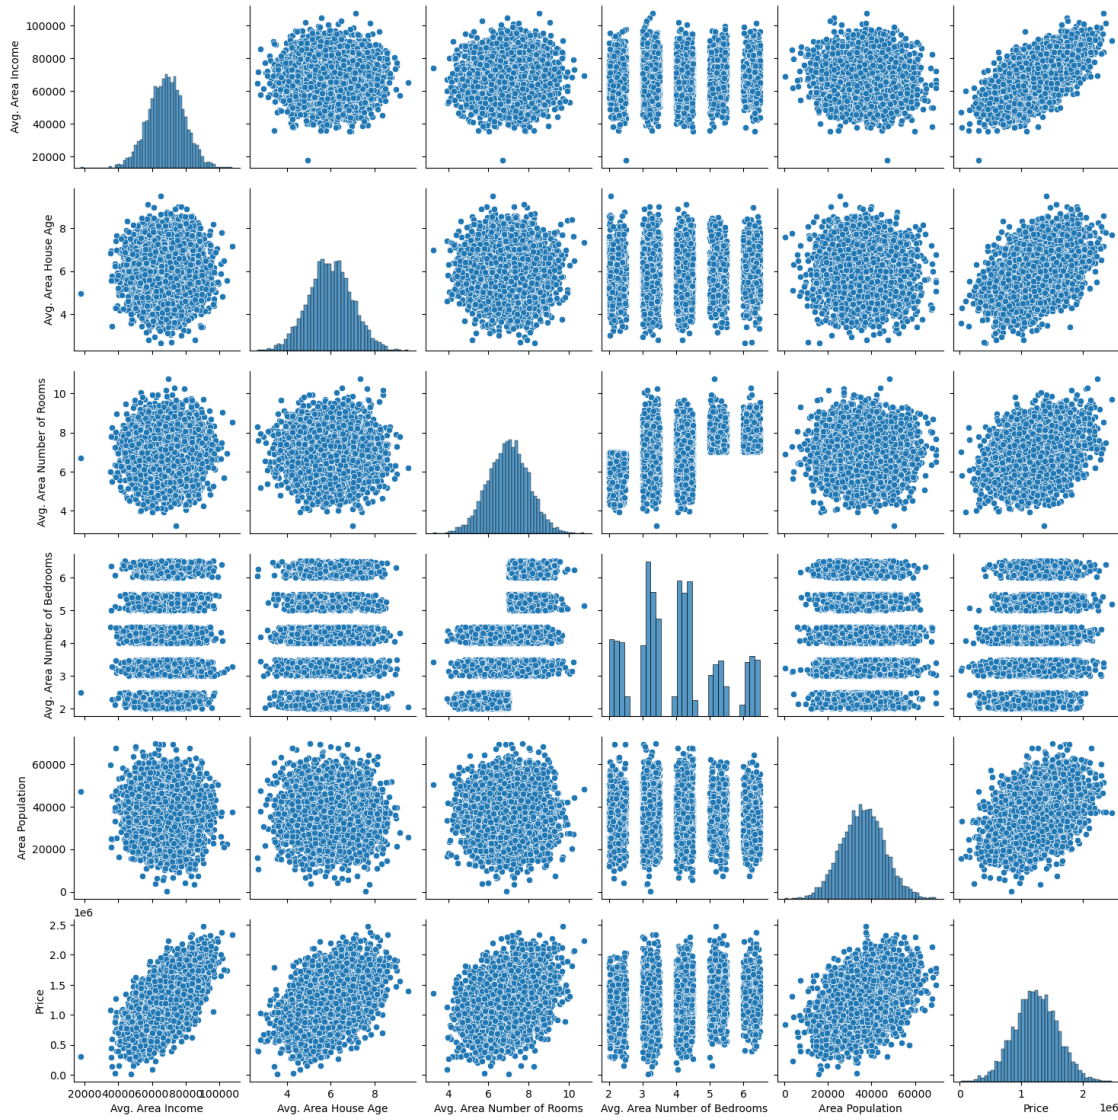
```
[24]: sns.boxplot(casas, showliers=False)
```

```
[24]: <AxesSubplot: >
```




```
[25]: sns.pairplot(casas)
```

```
[25]: <seaborn.axisgrid.PairGrid at 0x7f36990ad250>
```



```
[26]: sns.heatmap(casas.corr(numeric_only = True))
```

```
[26]: <AxesSubplot: >
```

```
[27]: casas.head(2)
```

```
[27]: Avg. Area Income Avg. Area House Age Avg. Area Number of Rooms \
0      79545.458574          5.682861          7.009188
1      79248.642455          6.002900          6.730821

Avg. Area Number of Bedrooms Area Population      Price \
0              4.09      23086.800503  1.059034e+06
1              3.09      40173.072174  1.505891e+06

Address
0  208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1  188 Johnson Views Suite 079\nLake Kathleen, CA...
```

```
[28]: X = casas[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of
Rooms',
              'Avg. Area Number of Bedrooms', 'Area Population']]
```

```
Y = casas['Price']
```

```
[29]: from sklearn.model_selection import train_test_split
```

```
X_treino, X_teste, y_treino, y_teste = train_test_split(X, Y, train_size=0.7,  
↪random_state=50)
```

```
[30]: from sklearn.linear_model import LinearRegression
```

```
[31]: LR = LinearRegression()
```


```
[32]: LR.fit(X_treino, y_treino)
```

```
[32]: LinearRegression()
```

```
[33]: predição = LR.predict(X_teste)
```

```
[34]: plt.scatter(y_teste, predição)
```

```
[34]: <matplotlib.collections.PathCollection at 0x7f372c6e44d0>
```



```
[35]: from sklearn.metrics import mean_absolute_error, mean_squared_error

mae = mean_absolute_error(y_true=y_teste, y_pred=predição)
#squared True returns MSE value, False returns RMSE value.
mse = mean_squared_error(y_true=y_teste, y_pred=predição) #default=True
rmse = mean_squared_error(y_true=y_teste, y_pred=predição, squared=False)

print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
```

```
MAE: 80728.93384538879
MSE: 10077066685.864893
RMSE: 100384.59386711137
```

0.4 Exercício

Crie e treine um modelo de regressão linear para prever a quantidade de banheiros que corresponde à variável alvo (Avg. Area Number of Bedrooms), utilizando a quantidade de quartos presentes nas casas (Avg. Area Number of Rooms) da base de dados USA_Housing.csv.

```
[ ]:
```