

Regressão_Logística_em_Python

April 29, 2025

1 Regressão Logística em Python

2 Importar as bibliotecas

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

3 Inserir os dados

```
[ ]: titanic = pd.read_csv('titanic_train.csv')
```

```
[ ]: titanic.head()
```

```
[ ]: 
```

	PassengerId	Survived	Pclass	...	Fare	Cabin	Embarked
0	1	0	3	...	7.2500	NaN	S
1	2	1	1	...	71.2833	C85	C
2	3	1	3	...	7.9250	NaN	S
3	4	1	1	...	53.1000	C123	S
4	5	0	3	...	8.0500	NaN	S

[5 rows x 12 columns]

```
[ ]: titanic.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Sex             891 non-null   object
```

```

5   Age          714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

```

```
[ ]: titanic.isnull()
```


```
[ ]:
   PassengerId  Survived  Pclass   Name  ... Ticket   Fare  Cabin  Embarked
0            False     False   False  False ... False  False   True    False
1            False     False   False  False ... False  False  False    False
2            False     False   False  False ... False  False   True    False
3            False     False   False  False ... False  False  False    False
4            False     False   False  False ... False  False   True    False
..           ...      ...      ...      ... ..  ...      ...
886          False     False   False  False ... False  False   True    False
887          False     False   False  False ... False  False  False    False
888          False     False   False  False ... False  False   True    False
889          False     False   False  False ... False  False  False    False
890          False     False   False  False ... False  False   True    False

```

```
[891 rows x 12 columns]
```


```
[ ]: sns.heatmap(titanic.isnull(), yticklabels=False, cbar=False)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f74df41ca50>
```




```
[ ]: sns.countplot(x='Survived', hue='Sex', data=titanic)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f74d685f350>
```




```
[ ]: sns.countplot(x='Survived', hue='Pclass', data=titanic)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f74df41c0d0>
```




```
[ ]: sns.countplot(x='SibSp', data=titanic)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f74d4f84910>
```




```
[ ]: sns.displot(titanic['Age'].dropna())
```

```
[ ]: <seaborn.axisgrid.FacetGrid at 0x7f74d508f990>
```




```
[ ]: sns.displot(titanic[titanic['Fare']<50])
```

```
[ ]: <seaborn.axisgrid.FacetGrid at 0x7f74df0fd750>
```



```
[ ]: sns.boxplot(x='Pclass', y='Age', data=titanic)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f74d4d29650>
```



```
[ ]: def input_idade(cols):
    Age = cols[0]
    Pclass = cols[1]


    if pd.isnull(Age):
        if Pclass == 1:
            return 39
        elif Pclass == 2:
            return 30
        else:
            return 25

    else:
        return Age
```

```
[ ]: titanic['Age'] = titanic[['Age', 'Pclass']].apply(input_idade, axis=1)
```

```
[ ]: sns.heatmap(titanic.isnull(), yticklabels=False, cbar=False)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f74df453610>
```

```
[ ]: titanic.drop('Cabin', axis=1, inplace=True)
```

```
[ ]: titanic.head()
```

```
[ ]:
  PassengerId  Survived  Pclass  ...    Ticket     Fare  Embarked
0            1         0       3  ...    A/5 21171    7.2500        S
1            2         1       1  ...    PC 17599   71.2833        C
2            3         1       3  ...  STON/O2. 3101282    7.9250        S
3            4         1       1  ...    113803   53.1000        S
4            5         0       3  ...    373450    8.0500        S
```

[5 rows x 11 columns]

```
[ ]: sexo=pd.get_dummies(titanic['Sex'],drop_first=True)
```

```
[ ]: embarque=pd.get_dummies(titanic['Embarked'],drop_first=True)
```

```
[ ]: titanic.drop(['Sex', 'Name', 'Ticket', 'Embarked'], axis=1, inplace=True)
```

```
[ ]: titanic.head()
```

```
[ ]: PassengerId  Survived  Pclass   Age  SibSp  Parch    Fare
0         1         0         3  22.0     1     0   7.2500
1         2         1         1  38.0     1     0  71.2833
2         3         1         3  26.0     0     0   7.9250
3         4         1         1  35.0     1     0  53.1000
4         5         0         3  35.0     0     0   8.0500
```

```
[ ]: titanic = pd.concat([titanic,sexo,embarque], axis=1)
```

```
[ ]: titanic.head()
```

```
[ ]: PassengerId  Survived  Pclass   Age  SibSp  Parch    Fare  male  Q  S
0         1         0         3  22.0     1     0   7.2500    1  0  1
1         2         1         1  38.0     1     0  71.2833    0  0  0
2         3         1         3  26.0     0     0   7.9250    0  0  1
3         4         1         1  35.0     1     0  53.1000    0  0  1
4         5         0         3  35.0     0     0   8.0500    1  0  1
```

#Regressão Logística

```
[ ]: from sklearn.model_selection import train_test_split
```

```
[ ]: X_treino, X_teste, Y_treino, Y_teste = train_test_split(titanic.
↳ drop('Survived',axis=1), titanic['Survived'], train_size=0.7,↳
↳ random_state=56)
```

```
[ ]: from sklearn.linear_model import LogisticRegression
```

```
[ ]: RL=LogisticRegression()
```

```
[ ]: RL.fit(X_treino,Y_treino)
```

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:940:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

```
[ ]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
warm_start=False)
```

```
[ ]: predictor = RL.predict(X_teste)
```

```
[ ]: from sklearn.metrics import classification_report  
     from sklearn.metrics import confusion_matrix
```

```
[ ]: print(classification_report(Y_teste, predictor))
```

	precision	recall	f1-score	support
0	0.86	0.89	0.88	178
1	0.76	0.72	0.74	90
accuracy			0.83	268
macro avg	0.81	0.80	0.81	268
weighted avg	0.83	0.83	0.83	268

```
[ ]: print(confusion_matrix(Y_teste, predictor))
```

```
[[158  20]  
 [ 25  65]]
```